


Spatial Data Management

Relational Representation of Spatial Data

- *Example:* Representation of geometric objects (here: fields of land) in normalized relations. Redundancy free representation requires distribution of information over 3 tables!



| Fields | |
|----------------|-----------------|
| FID | BID |
| F ₁ | B ₁ |
| F ₁ | B ₂ |
| F ₁ | B ₃ |
| F ₁ | B ₄ |
| F ₄ | B ₂ |
| F ₄ | B ₅ |
| F ₄ | B ₆ |
| F ₄ | B ₇ |
| F ₄ | B ₈ |
| F ₄ | B ₉ |
| F ₇ | B ₁₀ |
| F ₇ | B ₁₁ |
| F ₇ | B ₁₂ |

| Borders | | |
|-----------------|------------------|------------------|
| BID | PID ₁ | PID ₂ |
| B ₁ | P ₁ | P ₂ |
| B ₂ | P ₂ | P ₃ |
| B ₃ | P ₃ | P ₄ |
| B ₄ | P ₄ | P ₁ |
| B ₅ | P ₂ | P ₅ |
| B ₆ | P ₅ | P ₆ |
| B ₇ | P ₆ | P ₇ |
| B ₈ | P ₇ | P ₈ |
| B ₉ | P ₈ | P ₃ |
| B ₁₀ | P ₆ | P ₉ |
| B ₁₁ | P ₉ | P ₁₀ |
| B ₁₂ | P ₁₀ | P ₇ |

| Points | | |
|-----------------|------------------|------------------|
| PID | XCoord | YCoord |
| P ₁ | X _{P1} | Y _{P1} |
| P ₂ | X _{P2} | Y _{P2} |
| P ₃ | X _{P3} | Y _{P3} |
| P ₄ | X _{P4} | Y _{P4} |
| P ₅ | X _{P5} | Y _{P5} |
| P ₆ | X _{P6} | Y _{P6} |
| P ₇ | X _{P7} | Y _{P7} |
| P ₈ | X _{P8} | Y _{P8} |
| P ₉ | X _{P9} | Y _{P9} |
| P ₁₀ | X _{P10} | Y _{P10} |

Relational Representation of Spatial Data

- For (spatial) queries involving fields it is necessary to reconstruct the spatial information from the different tables
- E.g.: if we want to determine if a given point P is inside field F_2 , we have to find all corner-points of parcel F_2 first

```
SELECT Points.PID, XCoord, YCoord
FROM Fields, Border, Points
WHERE FID = 'F2' AND
Fields.BID = Borders.BID AND
(Borders.PID1 = Points.PID OR
Borders.PID2 = Points.PID)
```

- Even this simple query requires expensive joins of 3 tables
- Querying the geometry is not directly supported (P in F_2 ?)


Extension of Relational Model for Spatial Data

- Integration of spatial data types and operations into the core of a DBMS (→ object-oriented and object-relational databases)
 - Data types such as *Point*, *Line*, *Polygon*
 - Operations such as *ObjectIntersect*, *RangeQuery*, etc.
- Advantages:
 - Natural extension of the relational model and query languages
 - Facilitates design and querying of spatial databases
 - Spatial data types and operations can be supported by spatial index structures and efficient algorithms, implemented in the core of a DBMS
- All major database vendors today implement support for spatial data and operations in their database systems via object-relational extensions

Extension of Relational Model for Spatial Data

■ Example:

ForestZones(Zone:*Polygon*,ForestOfficial:*String*,Area:*Cardinal*)



| ForestZones | | |
|----------------|----------------|------------------------|
| Zone | ForestOfficial | Area (m ²) |
| R ₁ | Stevens | 3900 |
| R ₂ | Behrens | 4250 |
| R ₃ | Lee | 6700 |
| R ₄ | Goebel | 5400 |
| R ₅ | Jones | 1900 |
| R ₆ | Kent | 4600 |


- The province decides that a reforestation is necessary in an area described by a polygon S. Find all forest officials affected by this decision.

```
SELECT ForestOfficial
FROM ForestZones
WHERE ObjectIntersects (S,Zone)
```


Data Types for Spatial Objects

- Spatial objects are described by
 - Spatial Extent
 - *location* and/or *boundary* with respect to a reference point in a coordinate system, which is at least 2-dim.
 - Basic object types: *Point*, *Lines*, *Polygon*
 - Other Non-Spatial Attributes
 - Thematic attributes such as height, area, name, etc.


2-dim. points



2-dim. lines



2-dim. polygons




Spatial Query Processing

- DBMS has to support two types of operations
 - Operations to retrieve certain subsets of spatial object from the database
 - “Spatial Queries/Selections”, e.g., window query, point query, etc.
 - Operations that perform basic geometric computations and tests
 - E.g., point in polygon test, intersection of two polygons etc.
- Spatial selections, e.g. in geographic information systems, are often supported by an interactive graphical user interface

Basic Spatial Queries


- Containment Query: Given a spatial object R, find all objects that completely contain R. If R is a point: Point Query
- Region Query: Given a region R (polygon or circle), find all spatial objects that intersect with R. If R is a rectangle: Window Query
- Enclosure Query: Given a polygon region R, find all objects that are completely contained in R
- K-Nearest Neighbor Query: Given an object P, find the k objects that are closest to P (typically for points)




Containment Query



Point Query




Region Query



Window Query



Enclosure Query



2-nn Query

Basic Spatial Queries - Spatial Join

- Given two sets of spatial objects (typically minimum bounding rectangles)


$$S_1 = \{R_1, R_2, \dots, R_m\} \text{ and } S_2 = \{R'_1, R'_2, \dots, R'_n\}$$

- Spatial Join: Compute all pairs of objects (R, R') such that

$$R \in S_1, R' \in S_2 \text{ and } R \text{ intersects } R' (R \cap R' \neq \emptyset)$$

- Spatial predicates other than intersection are also possible, e.g. all pairs of objects that are within a certain distance from each other

$$\{A_1, A_2, \dots, A_6\} \otimes \{B_1, \dots, B_3\}$$



Spatial Join

Answer set

(A_5, B_1)

(A_4, B_1)

(A_1, B_2)


(A_6, B_2)

(A_2, B_3)

Index Support for Spatial Queries

- Conventional index structures such as B-trees are not designed to support spatial queries
 - Group objects only along one dimension
 - Do not preserve spatial proximity
- E.g. nearest neighbor query:



Nearest neighbor of Q is typically not the nearest neighbor in any single dimension



A and B closer in the X dimension;
C and D closer in the Y dimension.

Index Support for Spatial Queries (cont)

- Spatial index structures try to preserve spatial proximity
 - Group objects that are close to each other on the same data page
 - Problem: the number of bytes to store extended spatial objects (lines, polygons) varies
- Solution:
 - Store *Approximations* of spatial objects in the index structure, typically axis-parallel minimum bounding rectangles (MBR)
 - Exact object representation (ER) stored separately; pointers to ER in the index



Query Processing Using Approximations

■ Two-step procedure:

1. Filter Step:


- Use the index to find all approximations that satisfy the query
- Some objects already satisfy the query based on the approximation, others have to be checked in the refinement step → *Candidate Set*

2. Refinement Step:

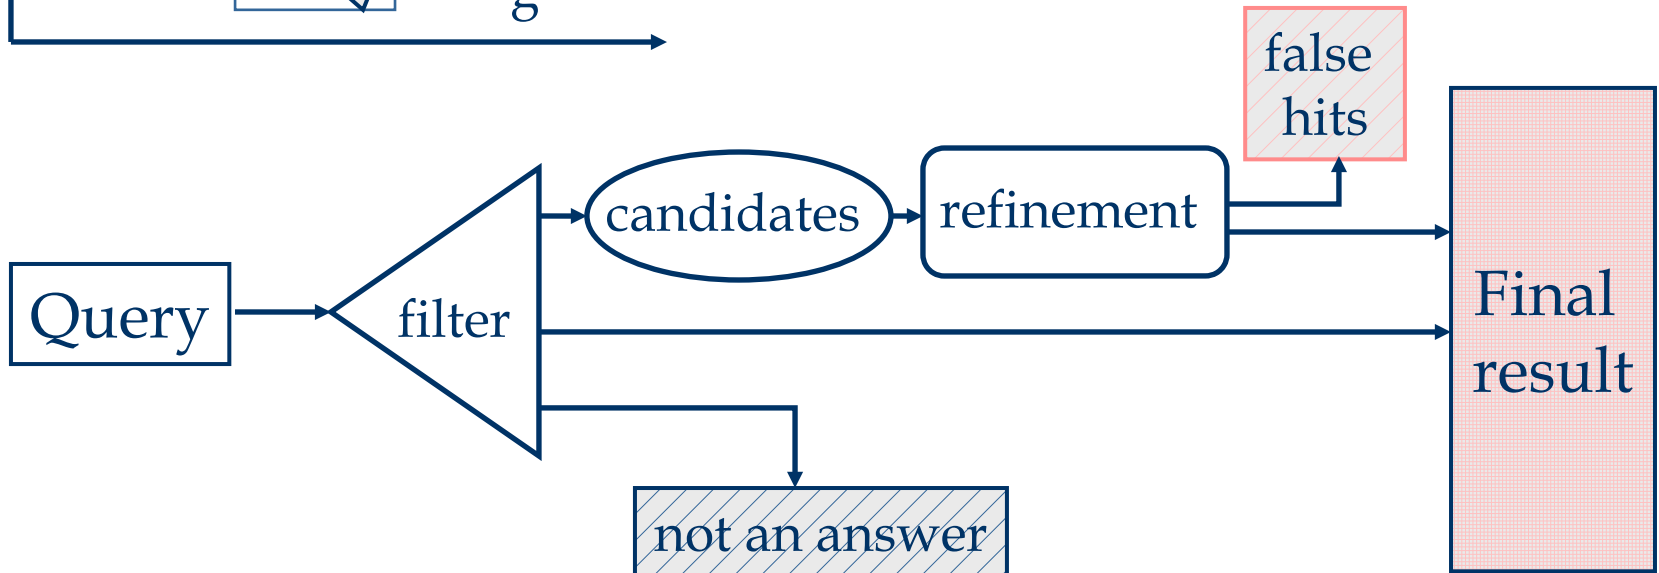
- Load the exact object representations for candidates left after the filter step and test whether they satisfies the query

Query Processing Using Approximations

Query window



- a and b are certain answers
- f, d, g are certainly not answers
- c and e are candidates
- c is a false hit



Embedding of the 2-dim. space into a 1-dim space


Basic Idea:

- The data space is partitioned into rectangular cells.
- Use a space filling curve to assign cell numbers to the cells (define a linear order on the cells)
 - The curve should preserve spatial proximity as good as possible
 - Cell numbers should be easy to compute
- Objects are approximated by cells.
- Store the cell numbers for objects in a conventional index structure with respect to the linear order

| | | | | | | | |
|----|----|----|----|----|----|----|----|
| 21 | 23 | 29 | 31 | 53 | 55 | 61 | 63 |
| 20 | 22 | 28 | 30 | 52 | 54 | 60 | 62 |
| 17 | 19 | 25 | 27 | 49 | 51 | 57 | 59 |
| 16 | 18 | 24 | 26 | 48 | 50 | 56 | 58 |
| 5 | 7 | 13 | 15 | 37 | 39 | 45 | 47 |
| 4 | 6 | 12 | 14 | 36 | 38 | 44 | 46 |
| 1 | 3 | 9 | 11 | 33 | 35 | 41 | 43 |
| 0 | 2 | 8 | 10 | 32 | 34 | 40 | 42 |

Space Filling Curves


■ Lexicographic order



| | | | | | | | |
|----|----|----|----|----|----|----|----|
| 56 | 57 | 58 | 59 | 60 | 61 | 62 | 63 |
| 48 | 49 | 50 | 51 | 52 | 53 | 54 | 55 |
| 40 | 41 | 42 | 43 | 44 | 45 | 46 | 47 |
| 32 | 33 | 34 | 35 | 36 | 37 | 38 | 39 |
| 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 |
| 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 |
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

Space Filling Curves (cont)


■ Hilbert Curve



| | | | | | | | |
|----|----|----|----|----|----|----|----|
| 63 | 62 | 49 | 48 | 47 | 44 | 43 | 42 |
| 60 | 61 | 50 | 51 | 46 | 45 | 40 | 41 |
| 59 | 56 | 55 | 52 | 33 | 34 | 39 | 38 |
| 58 | 57 | 54 | 53 | 32 | 35 | 36 | 37 |
| 5 | 6 | 9 | 10 | 31 | 28 | 27 | 26 |
| 4 | 7 | 8 | 11 | 30 | 29 | 24 | 25 |
| 3 | 2 | 13 | 12 | 17 | 18 | 23 | 22 |
| 0 | 1 | 14 | 15 | 16 | 19 | 20 | 21 |

Space Filling Curves

■ Z-Order



| | | | | | | | |
|----|----|----|----|----|----|----|----|
| 21 | 23 | 29 | 31 | 53 | 55 | 61 | 63 |
| 20 | 22 | 28 | 30 | 52 | 54 | 60 | 62 |
| 17 | 19 | 25 | 27 | 49 | 51 | 57 | 59 |
| 16 | 18 | 24 | 26 | 48 | 50 | 56 | 58 |
| 5 | 7 | 13 | 15 | 37 | 39 | 45 | 47 |
| 4 | 6 | 12 | 14 | 36 | 38 | 44 | 46 |
| 1 | 3 | 9 | 11 | 33 | 35 | 41 | 43 |
| 0 | 2 | 8 | 10 | 32 | 34 | 40 | 42 |

- Z-Order preserves spatial proximity relatively good
- Z-Order is easy to compute

Z-Order – Z-Values

■ Coding of Cells


- Partition the data space recursively into two halves
- Alternate X and Y dimension
- Left/bottom $\rightarrow 0$
- Right/top $\rightarrow 1$

■ Z-Value: (c, l)

c = decimal value of the bit string

l = level (number of bits)

- if all cells are on the same level, then l can be omitted



Z-Order - Representation of Spatial Objects


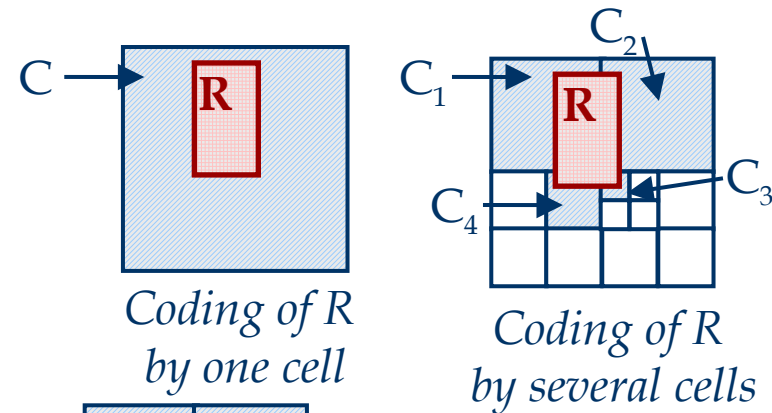
■ For Points

- Use a fixed a resolution of the space in both dimensions, i.e., each cell has the same size
- Each point is then approximated by one cell

| | | | | | | | |
|----|----|----|----|----|----|----|----|
| 21 | 23 | 29 | 31 | 53 | 55 | 61 | 63 |
| 20 | 22 | 28 | 30 | 52 | 54 | 60 | 62 |
| 17 | 19 | 25 | 27 | 49 | 51 | 57 | 59 |
| 16 | 18 | 24 | 26 | 48 | 50 | 56 | 58 |
| 5 | 7 | 13 | 15 | 37 | 39 | 45 | 47 |
| 4 | 6 | 12 | 14 | 36 | 38 | 44 | 46 |
| 1 | 3 | 9 | 11 | 33 | 35 | 41 | 43 |
| 0 | 2 | 8 | 10 | 32 | 34 | 40 | 42 |

■ For extended spatial object

- minimum enclosing cell
 - Problems with cells that intersect the first partitions already
- improvement: use several cells
 - Better approximation of the objects
 - Redundant storage
 - Redundant retrieval in spatial queries



Z-Order – Mapping to a B⁺-Tree

- Linear order for Z-values to store them in a B⁺-tree:

Let (c_1, l_1) and (c_2, l_2) be two Z-Values and let
 $l = \min\{l_1, l_2\}$.

The order relation \leq_Z (that defines a linear order on Z-values) is then defined by:

$$(c_1, l_1) \leq_Z (c_2, l_2) \text{ iff } (c_1 \text{ div } 2^{(l_1-l)}) \leq (c_2 \text{ div } 2^{(l_2-l)})$$


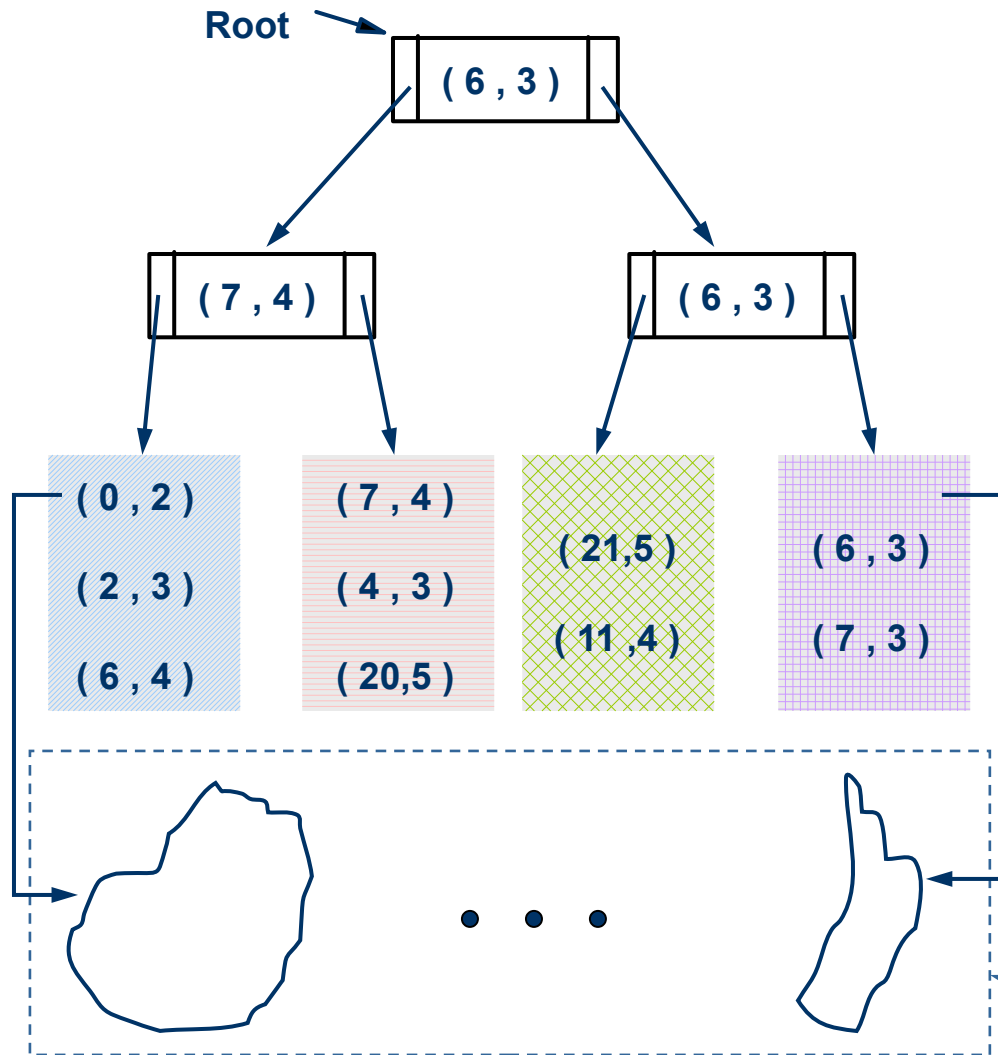
Examples:

$$(1,2) \leq_Z (3,2),$$

$$(3,4) \leq_Z (3,2),$$

$$(1,2) \leq_Z (10,4)$$

Mapping to a B⁺-Tree - Example





$(0, 2) \leq (7, 4) \leq (7, 4) \leq (6, 3)$
 $(2, 3) \leq (7, 4) \leq (4, 3) \leq (6, 3)$
 $(6, 4) \leq (7, 4) \leq (20, 5) \leq (6, 3)$

Exact representations stored in a different location

Mapping to a B⁺-Tree - Window Query


- Window Query → Range Query in the B⁺-tree
 - find all entries (Z-Values) in the range $[l, u]$ where
 - l = smallest Z-Value of the window (bottom left corner)
 - u = largest Z-Value of the window (top right corner)
 - l and u are computed with respect to the maximum resolution/length of the Z-values in the tree (here: 6)



Window: Min = (0,6), Max = (10,6)

R-Trees


- The R-tree is a tree-structured index that remains balanced on inserts and deletes.
- Each key stored in a leaf entry is intuitively a **box**, or collection of **intervals**, with one interval per dimension.
- Example in 2-D:




R-Tree Properties

- *Leaf entry* = $\langle \text{n-dimensional box, rid} \rangle$
 - *key value* being a box.
 - Box is the tightest bounding box for a data object.
- *Non-leaf entry* = $\langle \text{n-dim box, ptr to child node} \rangle$
 - Box covers all boxes in child node (in fact, subtree).
- All leaves at same distance from root.
- Nodes must be kept at least 50% full (except root).
 - Can choose a parameter m that is $\leq 50\%$, and ensure that every node is at least $m\%$ full.

Example of an R-Tree



Example R-Tree (cont.)



Search for Objects Overlapping **Box Q**

Start at **root**.

1. If current node is non-leaf, for each entry $\langle E, \text{ptr} \rangle$, if **box E** overlaps **Q**, search subtree identified by **ptr**.
2. If current node is leaf, for each entry $\langle E, \text{rid} \rangle$, if **E** overlaps **Q**, **rid** identifies an object that might overlap **Q**.

*Note: May have to search **several** subtrees at each node!
(In contrast, a B-tree equality search goes to just one leaf.)*

Improving Search Using Constraints


- It is convenient to **store boxes** in the R-tree as approximations of arbitrary regions, because boxes can be represented compactly.
- But why not use **convex polygons to approximate query regions** more accurately?
 - Will reduce overlap with nodes in tree, and reduce the number of nodes fetched by avoiding some branches altogether.
 - Cost of overlap test is higher than bounding box intersection, but it is a main-memory cost, and can actually be done quite efficiently. Generally a win.

Insert Entry $\langle B, ptr \rangle$

- Start at root and go down to “best-fit” leaf L .
 - Go to child whose box needs least enlargement to cover B ; resolve ties by going to smallest area child.
- If best-fit leaf L has space, insert entry and stop. Otherwise, split L into L_1 and L_2 .
 - Adjust entry for L in its parent so that the box now covers (only) L_1 .
 - Add an entry (in the parent node of L) for L_2 . (This could cause the parent node to recursively split.)

Splitting a Node During Insertion

- The entries in node L plus the newly inserted entry must be distributed between L_1 and L_2 .
- Goal is to reduce likelihood of both L_1 and L_2 being searched on subsequent queries.
- **Idea:** Redistribute so as to **minimize area** of L_1 plus area of L_2 .



R-Tree Variants

- The **R* tree** uses the concept of **forced reinserts** to reduce overlap in tree nodes. When a node overflows, instead of splitting:
 - Remove some (say, 30% of the) entries and reinsert them into the tree.
 - Could result in all reinserted entries fitting on some existing pages, avoiding a split.
- R* trees also use a different heuristic, minimizing **box perimeters** rather than **box areas** during insertion.
- Another variant, the **R+ tree**, avoids overlap by inserting an object into multiple leaves if necessary.
 - Searches now take a single path to a leaf, at cost of redundancy.

GiST

- The Generalized Search Tree (GiST) abstracts the “tree” nature of a class of indexes including B+ trees and R-tree variants.
- Striking similarities in insert/delete/search and even concurrency control algorithms make it possible to provide “templates” for these algorithms that can be customized to obtain the many different tree index structures.
- B+ trees are so important (and simple enough to allow further specialization) that they are implemented specially in all DBMSs.
- GiST provides an alternative for implementing other tree indexes in an ORDBS.

Indexing High-Dimensional Data

- Typically, high-dimensional datasets are collections of points, not regions.
 - E.g., Feature vectors in multimedia applications.
 - Very sparse
- Nearest neighbor queries are common.
 - R-tree becomes worse than sequential scan for most datasets with more than a dozen dimensions.
- As dimensionality increases **contrast** (ratio of dist. between nearest and farthest points) usually decreases; “nearest neighbor” is not meaningful.
 - In any given data set, advisable to empirically test contrast.

Summary

- Spatial data management has many applications, including GIS, CAD/CAM, multimedia indexing.
 - Point and region data
 - Overlap/containment and nearest-neighbor queries
- Many approaches to indexing spatial data
 - R-tree approach is widely used in GIS systems
 - Other approaches include Grid Files, Quad trees, and techniques based on “space-filling” curves.
 - For high-dimensional datasets, unless data has good “contrast”, nearest-neighbor may not be well-separated

Comments on R-Trees

- Deletion consists of searching for the entry to be deleted, removing it, and if the node becomes under-full, deleting the node and then re-inserting the remaining entries.
- Overall, works quite well for 2 and 3 D datasets. Several variants (notably, R+ and R* trees) have been proposed; widely used.
- Can improve search performance by using a convex polygon to approximate query shape (instead of a bounding box) and testing for polygon-box intersection.